

An Vo

(+1) 734-649-1118 | voan@umich.edu | vokhanhan25@gmail.com

[in](#) [vokhanhan25](#) | [anvo25](#) | [an_vo12](#) | [Google Scholar](#) | [anvo.me](#)

PhD student in Computer Science & Engineering @ University of Michigan
3944 BBB, University of Michigan, Ann Arbor, MI 48109, USA

BIOGRAPHY

An is a PhD student in Computer Science and Engineering at the [University of Michigan](#), advised by [Rada Mihalcea](#). His research focuses broadly on foundation models (LLMs/VLMs), with an emphasis on making them more trustworthy and explainable. His research has been [used in industry](#) by Google DeepMind, Alibaba Qwen, ByteDance, and Moonshot AI, and [featured on Sky News](#). He also collaborates closely with [Anh Totti Nguyen](#). Before joining Michigan, he worked as a Research Engineer at [MBZUAI](#) with [Thamar Solorio](#) and earned his MS from [KAIST](#), where he received the Outstanding Thesis Award and was fully funded by the [Hyundai CMK Global Scholarship](#).

EDUCATION

- **University of Michigan** Aug 2026 – May 2031 (expected)
PhD in Computer Science and Engineering | Advisor: [Rada Mihalcea](#) Ann Arbor, US
- **Korea Advanced Institute of Science & Technology (KAIST)** Feb 2024 – Feb 2026
M.S. in Computer Science | Advisor: [Daeyoung Kim](#) | Co-advisor: [Anh Totti Nguyen](#) Daejeon, South Korea
 - Thesis: Detecting and Understanding Biases in Foundation Models (*Outstanding Thesis Award*)
- **University of Information Technology (UIT), Vietnam National University HCMC** Aug 2019 – Apr 2023
B.S. in Computer Science | Advisor: [Ngoc Hoang Luong](#) Ho Chi Minh City, Vietnam
 - GPA: 9.32/10 (*Valedictorian, #1*)
 - Thesis: Many-Objective Evolutionary Neural Architecture Search with Performance Predictors (10/10–Best Thesis)

EXPERIENCE

- **University of Michigan** Aug 2026 – Present
Graduate Student Research Assistant (GSRA) Ann Arbor, US
 - Working on frontier ML/NLP problems
- **Mohamed bin Zayed University of Artificial Intelligence (MBZUAI)** Jan 2026 – Jul 2026
Research Enginner Abu Dhabi, UAE
 - Working on Multimodal LLMs, NLP with [Thamar Solorio](#)
 - Supporting the different research projects in the lab, contributing to research discussions and writing papers.
 - Engaging with the larger NLP and MBZUAI community
- **Korea Advanced Institute of Science & Technology (KAIST)** Feb 2024 – Jan 2025
Graduate Research Assistant Daejeon, South Korea
 - Broadly interested (but not limited to) in foundation models (e.g, LLMs, VLMs), especially in making them more trustworthy and explainable in rare/hard cases.
- **Auburn University** Mar 2024 – Jan 2025
Research Collaborator/Remote Visiting Student Remote
 - Remote collaboration with [Anh Totti Nguyen](#).
- **University of Information Technology (UIT), Vietnam National University HCMC** Apr 2020 – Jan 2024
Undergraduate Research Assistant Ho Chi Minh City, Vietnam
 - Conducted research on Evolutionary Computation, Neural Architecture Search (NAS), Vehicle Routing Problem, and Multi-Objective Optimization.
 - Developed novel algorithms to enhance the efficiency of NAS using evolutionary and other optimization algorithms.
- **ZaloPay, VNG Corporation** Jul 2023 – Sep 2023
Associate Data Scientist Ho Chi Minh City, Vietnam

- Led development of a Customer Lifetime Value (CLV) prediction system for 5M ZaloPay users, leveraging XGBoost/LightGBM and advanced financial feature engineering.
- Improved user segmentation and retention strategies by enabling more accurate value prediction and targeted engagement initiatives.

• Eximbank

Software Engineer

May 2023 – Jul 2023

Ho Chi Minh City, Vietnam

- Developed core banking modules to enhance system functionality and support business operations.
- Engineered database-ready features enabling efficient data extraction and analytics for the business team.

PUBLICATIONS

C=CONFERENCE, J=JOURNAL, P=PREPREPRINT

* indicates equal contribution.

Conference papers

- [C.1] An Vo*, Vy Tuong Dang*, Khai-Nguyen Nguyen, Emilio Villa-Cueva, Thamar Solorio†, Anh Totti Nguyent†, Daeyoung Kim† (2026). Deep and shallow biases in language models. The 2026 Conference on Empirical Methods in Natural Language Processing (EMNLP 2026). [[project website](#)]
- [C.2] Vy Tuong Dang*, An Vo*, Emilio Villa-Cueva, Quang Tau, Duc Dm, Thamar Solorio, Daeyoung Kim (2026). VMMU: A Vietnamese Multitask Multimodal Understanding and Reasoning Benchmark. *Findings of the Association for Computational Linguistics: ACL-IJCNLP 2026*. [[paper](#) | [project website](#)]
- [C.3] Thao Do, Dinh Phu Tran, An Vo, Seon Kwon Kim, Daeyoung Kim (2026). EnComp: Lightweight Encoder-Only Context Compression for Retrieval-Augmented Question Answering. *The 5th Asia-Pacific Chapter of the Association for Computational Linguistics & the 15th International Joint Conference on Natural Language Processing (ACL-IJCNLP 2026)*. [[paper](#)]
- [C.4] An Vo*, Khai-Nguyen Nguyen*, Mohammad Reza Taesiri, Vy Tuong Dang, Anh Totti Nguyent†, Daeyoung Kim† (2025). Vision Language Models are Biased. The Fourteenth International Conference on Learning Representations (ICLR 2026). [[paper](#) | [project website](#)]
- [C.5] Gabriel Manalu, An Vo, Duc Dm, Quang Tau, Nguyen Dinh Son, Minh Son Hoang, Hyeontaek Hwang, Jaesung Lee, Hyojin Choi, Tuan Phong Tran, Daeyoung Kim (2026). RADBench: Evaluating Foundation VLMs for Corner Case Description on A Comprehensive Dataset. *The IEEE Intelligent Vehicles Symposium (IV 2026)*. [[paper](#)]
- [C.6] An Vo, Mohammad Reza Taesiri, Daeyoung Kim, Anh Totti Nguyen (2025). B-score: Detecting biases in large language models using response history. *Forty-Second International Conference on Machine Learning (ICML 2025)*. [[paper](#) | [project website](#)]
- [C.7] Thao Do, Dinh Phu Tran, An Vo, Daeyoung Kim (2025). Reference-Based Post-OCR Processing with LLM for Precise Diacritic Text in Historical Document Recognition. *Annual AAAI Conference on Artificial Intelligence (AAAI 2025)*. [[paper](#) | [code](#) | [dataset](#)]
- [C.8] An Vo, Ngoc Hoang Luong (2024). Efficient Multi-Objective Neural Architecture Search via Pareto Dominance-based Novelty Search. *Genetic and Evolutionary Computation Conference (GECCO 2024)*. [[paper](#) | [code](#)]
- [C.9] Nhat Minh Le*, An Vo*, Ngoc Hoang Luong (2024). Zero-Cost Proxy-Based Hierarchical Initialization for Evolutionary Neural Architecture Search. *IEEE Congress on Evolutionary Computation (CEC 2024)*. [[paper](#)]
- [C.10] Khoa Huu Tran, Luc Truong, An Vo, Ngoc Hoang Luong (2023). Accelerating Gene-pool Optimal Mixing Evolutionary Algorithm for Neural Architecture Search with Synaptic Flow. *Genetic and Evolutionary Computation Companion Conference (GECCO Companion 2023)*. [[paper](#) | [code](#)]
- [C.11] An Vo, Tan Ngoc Pham, Van Bich Nguyen, Ngoc Hoang Luong (2022). Training-Free Multi-Objective and Many-Objective Evolutionary Neural Architecture Search with Synaptic Flow. *International Symposium on Information and Communication Technology (SoICT 2022)*. **Best Paper Award**. [[paper](#) | [code](#)]

Journal papers

- [J.1] An Vo, Nhat Minh Le, Ngoc Hoang Luong (2026). Lightweight Multi-Objective and Many-Objective Problem Formulations for Evolutionary Neural Architecture Search with the Training-Free Performance Metric Synaptic Flow. *SN Computer Science*. [[paper](#) | [code](#)]
- [J.2] Ngoc Hoang Luong, Quan Minh Phan, An Vo, Tan Ngoc Pham, Dzung Tri Bui (2024). Lightweight Multi-Objective Evolutionary Neural Architecture Search with Low-Cost Proxy Metrics. *Information Sciences*. [[paper](#) | [code](#)]
- [J.3] An Vo, Tan Ngoc Pham, Van Bich Nguyen, Ngoc Hoang Luong (2023). Efficient Multi-Objective Neural Architecture Search via Tree Search with Training-Free Metrics. *Informatica*. [[paper](#) | [code](#)]

SELECTED GRANTS AND FELLOWSHIPS

In PhD Program

- 2026: Rackham Conference Travel Funding (for EMNLP 2026) – \$1150

In M.S. Program

- 2025: Hyundai CMK Global Scholarship of Excellence (for ICML 2025 paper acceptance) – \$2200
- 2025: Hyundai CMK International Symposium Scholarship (ICML 2025 travel grant) – \$1500
- 2025: Cohere Labs Research Grant (with Vy Tuong Dang) – \$2,000
- 2025: Together AI Research Grant – \$150
- 2024: OpenAI Research Grant – \$2,500
- 2024: Cohere For AI Research Grant – \$1,000
- 2024–2026: KAIST Fellowship – \$22,000 (estimated total value)
- 2024–2026: Hyundai Chung-Mong Koo Global Scholarship – \$45,000 (estimated total value)

In B.S. Program

- 2022–2023: Incentive for Scientific Publications – \$100
- 2019–2023: University of Information Technology, VNU-HCM Merit Scholarships – \$3,600 (estimated total value)

In High School

- 2016–2019: Incentive for Participation and Awards in High School Olympiads – \$1,000 (estimated total value)

SELECTED HONORS AND AWARDS

In M.S. Program

- Outstanding M.S. Thesis Award

In B.S. Program

- 2023: **Valedictorian** of the B.S. Program
- 2023: Excellent Graduate of the B.S. Program
- 2022: **Best Paper Award** at the International Symposium on Information and Communication Technology (SoICT)
- 2022: Ho Chi Minh City Outstanding Youth
- 2022: Award for Outstanding Scientific Research Publications
- 2019: Top 100 Student Leaders of Ho Chi Minh City
- 2019–2023: Recognized for Excellence in Academics and Personal Development

In High School

- 2019: Consolation Prize in the Provincial Science and Engineering Fair
- 2018: First Prize in the Provincial Youth Informatics Competition
- 2018: First Prize in the Provincial English-Language Science and Engineering Competition in Informatics
- 2018–2019: Competed in the National Olympiad in Informatics (VOI) x2
- 2017–2018: Second Prize in the Provincial Olympiad in Informatics x2
- 2017: Bronze Medal in the Summer Olympiad in the Mekong Delta in Informatics
- 2017: Bronze Medal in the April 30th Olympiad in Informatics
- 2017: Consolation Prize in the Provincial Physics Olympiad via Internet

SELECTED PRESS COVERAGE

See details and links [here](#).

Industry Usage

- 2026: Moonshot AI – [Used VLMs are Biased as a source benchmark for developing PerceptionBench.](#)
- 2026: Alibaba Qwen – [Used VLMs are Biased to evaluate Qwen3.8.](#)
- 2026: ByteDance – [Used VLMs are Biased to evaluate Seed 2.0.](#)
- 2025: ByteDance – [Used VLMs are Biased to evaluate Seed 1.8.](#)
- 2025: Google DeepMind – [Used VLMs are Biased to evaluate Gemini 3 Pro.](#)
- 2025: Moondream – [Used examples from VLMs are Biased to demonstrate reduced counting bias through grounded reasoning.](#)

News Coverage

- 2026: [Vision Language Models are Biased](#) is featured in this Sky News [film](#) and [article](#) discussing the role of AI in modern warfare and its potential dangers.

Articles & Blog Posts

- 2026: KAIST School of Computing – [ICLR 2026 paper reveals memorization bias and limitations in counterfactual reasoning in vision-language models.](#)
- 2025: Gary Marcus – [GPT-5: Overdue, overhyped and underwhelming](#)
- 2025: Hacker News (Vision Language Models are Biased) – [Front page](#)

MENTORING

Undergraduate Students at University of Information Technology, Vietnam National University–Ho Chi Minh City

- [Khoa Huu Tran](#) (Oct 2022 – Jan 2024) – [GECCO Late-Breaking Abstract Paper](#)
- [Luc Truong](#) (Oct 2022 – Jan 2024) – [GECCO Late-Breaking Abstract Paper](#)
- [Minh Le](#) (Oct 2022 – Feb 2023) – [IEEE CEC Paper](#)
- [Vy Tuong Dang](#) (Dec 2021 – Aug 2025) – KAIST Scholarship for M.S. Program (Full Tuition Fee & Stipend) & [VMMU](#)

SELECTED PROFESSIONAL SERVICE

- **Conference Reviewer/Program Committee:** [NeurIPS 2026](#), [AAAI 2026](#), [BMVC 2025](#)
- **Journal Reviewer:** [IEEE Transactions on Evolutionary Computation](#)

TALKS

- 2024: **AutoID Labs Asia Workshop 2024** – Efficient Multi-Objective Neural Architecture Search via Pareto Dominance-based Novelty Search.
- 2022: **SoICT 2022** – Training-Free Multi-Objective and Many-Objective Evolutionary Neural Architecture Search with Synaptic Flow.

SELECTED CERTIFICATES

- 2025: IELTS 7.0 Overall (Exam date: 11 Oct 2025)
- 2023: TOEIC 880 (Exam date: 04 Mar 2023)
- 2021: Computational Thinking for Problem Solving – University of Pennsylvania, Coursera
- 2021: Linear Algebra for Machine Learning and Data Science – DeepLearning.AI, Coursera
- 2021: Python for Data Science and AI – IBM, Coursera
- 2021: Improving Deep Neural Networks: Hyperparameter Tuning, Regularization and Optimization – DeepLearning.AI, Coursera
- 2021: Neural Networks and Deep Learning – DeepLearning.AI, Coursera
- 2021: Fundamentals of Reinforcement Learning – University of Alberta, Coursera
- 2021: Sample-based Learning Methods – University of Alberta, Coursera

- 2020: How Google Does Machine Learning – Google Cloud, Coursera
- 2020: What is Data Science? – IBM, Coursera
- 2020: Basic Statistics – University of Amsterdam, Coursera
- 2020: Machine Learning – Andrew Ng, Stanford University, Coursera

REFERENCES

Dr. Rada Mihalcea

Michigan AI Director & Professor at UMich
mihalcea at umich.edu

Dr. Daeyoung Kim

Professor at KAIST
kimd at kaist.ac.kr

Dr. Thamar Solorio

Vice Provost & Professor at MBUAI
thamar.soloio at mbzuai.ac.ae

Dr. Anh Totti Nguyen

Associate Professor at Auburn University
anh.ng8 at gmail.com

Dr. Mohammad Reza Taesiri

Research Scientist at EA SPORTS
mtaesiri at gmail.com

Dr. Ngoc Hoang Luong

Head of AI Dept & Lecturer at UIT, VNU-HCM
hoangln at uit.edu.vn